

Driving Questions

What are the workarounds of transforming the images to make them of the same size, before using them to train the classification model? One such alternative is padding. What are its trade offs? Has it been used in similar problems?

1. Resizing Techniques

[Evaluating the Uncertainty of Classification Due to Image Resizing Techniques for Satellite Image Classification - Springer](#)

Investigates how different image resizing methods impact the performance and uncertainty of convolutional neural network (CNN) models in satellite image classification.

Objective/Methodology:

1. To assess how various image resizing techniques affect the accuracy and uncertainty of CNN-based satellite image classification.
2. Resizing Techniques Evaluated:
 - Nearest Neighbor Interpolation: Assigns the value of the nearest pixel without considering surrounding pixels, which can lead to jagged edges.
 - Bilinear Interpolation: Calculates the average of the four nearest pixels, resulting in smoother images.
 - Bicubic Interpolation: Uses the closest 16 pixels to compute a weighted average, producing even smoother images but at a higher computational cost.
3. Performance Metrics:
 - Accuracy: The proportion of correctly classified images.
 - Sensitivity (Recall): The ability to correctly identify positive instances.
 - Specificity: The ability to correctly identify negative instances.
 - Uncertainty: Measured using Shannon entropy to quantify the confidence of the model's predictions.

Findings:

- Accuracy: The model achieved the highest accuracy with bilinear interpolation, followed by bicubic, and the lowest with nearest neighbor interpolation.
- Uncertainty: Bilinear interpolation resulted in the lowest uncertainty (entropy), indicating more confident predictions, while nearest neighbor interpolation led to higher uncertainty.
- Sensitivity and Specificity: Both metrics were superior when using bilinear interpolation compared to the other methods.
- The study highlights that the choice of image resizing technique significantly influences the performance and confidence of CNN models in satellite image classification. **Bilinear interpolation emerged as the most effective method among those tested, balancing computational efficiency and classification accuracy.**

[Learning to Resize in Computer Vision - Keras](#)

The Keras article "Learning to Resize in Computer Vision" demonstrates the implementation of a learnable image resizing module, as proposed in the paper "Learning to Resize Images for Computer Vision Tasks" by Talebi and Milanfar (2021).

Objective/Methodology:

1. The primary goal is to enhance the performance of vision models by replacing traditional, fixed image resizing methods (e.g., bilinear or bicubic interpolation) with a learnable resizing module. This module is trained jointly with the main vision model to optimize the resizing process for the specific task at hand.
2. Learnable Resizer Module:
 - The module consists of convolutional layers followed by residual blocks, designed to learn optimal representations of images at a target resolution (e.g., 150×150 pixels).
 - It begins with a naive resizing step using bilinear interpolation, followed by a series of convolutional and residual blocks to refine the resized image.
3. Integration with Vision Model:
 - The learnable resizer is integrated into a DenseNet-121 architecture.
 - The combined model is trained end-to-end, allowing the resizer to adapt its parameters to improve the performance of the DenseNet classifier.

Findings:

The implementation demonstrates that incorporating the learnable resizer leads to improved classification accuracy compared to using traditional resizing methods.

Specifically, the model with the learnable resizer achieved a top-1 accuracy of 67.67%, while the model without it achieved 60.19%. By allowing the resizing process to be optimized during training, models can achieve better performance without significant increases in computational complexity. For a walkthrough and code implementation, refer to the Keras link.

2. Padding Techniques

[Position, Padding and Predictions: A Deeper Look at Position Information in CNNs - Springer](#)

The paper "Position, Padding and Predictions: A Deeper Look at Position Information in CNNs" investigates how padding strategies in Convolutional Neural Networks (CNNs) influence the encoding of absolute position information and affect model performance.

Findings:

1. Padding and Position Encoding:
 - Zero padding, commonly used to maintain spatial dimensions during convolution, inadvertently introduces absolute position information into CNNs.

- This positional encoding arises because zero padding creates distinct boundary effects, allowing the network to infer spatial locations within the input image.
- 2. Impact on Semantic Representations:
 - The study reveals that the presence of position information can both enhance and impair the learning of semantic features, depending on the task.
 - For tasks where spatial context is crucial, such as object detection, position encoding can be beneficial. Conversely, for tasks requiring spatial invariance, like certain classification problems, it may introduce biases.
- 3. Alternative Padding Strategies:
 - The authors compare zero padding with other methods, including reflection and replication padding.
 - They find that while zero padding encodes the most position information, alternative strategies can mitigate unwanted positional biases, leading to improved performance in specific applications.

Implications:

In our work comparing images with different stains, it's crucial to consider how padding affects the model's ability to recognize spatial patterns consistently. There are many different padding methods, and choosing one that aligns with the spatial characteristics of our data can enhance the model's performance in identifying similarities between samples. Yet, padding can introduce positional information, which might bias a model by making boundary areas distinctive in ways that could affect model generalization.

[The Impact of Padding on Image Classification by Using Pre-trained Convolutional Neural Networks - Springer](#)

In the study "The Impact of Padding on Image Classification by Using Pre-trained Convolutional Neural Networks," researchers investigated how different padding strategies affect the performance of pre-trained convolutional neural networks (CNNs) in image classification tasks.

Findings:

1. Padding Strategies Evaluated:
 - Constant Padding: Surrounding the image with a constant value, typically zero.
 - Edge Padding: Extending the edges of the image outward.
 - Mirror Padding: Reflecting the image content along its borders.
2. Performance Metrics:
 - The study assessed Top-1 and Top-5 accuracy, as well as the average confidence of the classifier.
3. Results:
 - For low-resolution images, mirror padding yielded the highest classification accuracy among the strategies tested.
 - Edge padding also improved performance but was less effective than mirror padding.

- Constant padding resulted in the lowest accuracy, likely due to the introduction of artificial borders that do not represent actual image content.

Implications:

In medical imaging, particularly with histopathological samples, images often vary in size and resolution. Applying appropriate padding strategies can enhance the performance of pre-trained CNNs by preserving essential features and minimizing artifacts introduced during preprocessing. The findings suggest that **mirror padding** is a preferable method when dealing with high-resolution medical images, as it maintains continuity at the image borders, leading to better classification outcomes.

3. Applications in Similar Problems

[Introducing Hann windows for reducing edge-effects in patch-based image segmentation - Plos Journals](#)

The study aims to address edge artifacts in patch-based image segmentation for large medical images. Specifically, it investigates the application of Hann window functions to reduce boundary effects in convolutional neural networks (CNNs) used for patch-based segmentation.

Methodology:

1. Patch-Based Segmentation: Large medical images are split into smaller patches due to computational limits. This approach often results in visible artifacts at patch boundaries, which can degrade segmentation accuracy.
2. Hann Window Application: The authors propose using a Hann window function, which applies a smooth weighting to pixel values, emphasizing the central pixels of each patch and reducing the influence of edge pixels.
3. Comparison and Evaluation: The performance of patch-based segmentation using Hann windows was compared to traditional segmentation techniques. Quantitative metrics and qualitative visual inspections were used to evaluate the results, including metrics such as Dice coefficient and Intersection over Union (IoU).

Findings:

- The use of Hann windows effectively minimized edge artifacts in the segmentation process, resulting in smoother transitions between patches and improving the visual quality of the segmented images.
- Quantitatively, segmentation with Hann windows outperformed conventional patch-based segmentation methods, showing higher accuracy and consistency, especially around patch boundaries.
- The authors demonstrated that the Hann window method could be seamlessly integrated into CNN workflows without requiring model retraining, making it a highly adaptable solution.

Best Practices for Preparing and Augmenting Image Data for CNNs - Machine Learning Mastery

The article "Best Practices for Preparing and Augmenting Image Data for Convolutional Neural Networks" by Jason Brownlee provides a comprehensive overview of effective techniques for preprocessing and augmenting image data to enhance the performance of Convolutional Neural Networks (CNNs).

Objective/Methodology:

1. Data Preparation:
 - Centering: Subtracting the per-channel mean pixel values calculated from the training dataset to center the data around zero.
 - Normalization: Scaling pixel values to a standard range, such as $[0, 1]$ or $[-1, 1]$, to facilitate faster convergence during training.
2. Training Data Augmentation:
 - Random Rescaling: Applying random zooms to images to make the model invariant to object size variations.
 - Horizontal Flips: Flipping images horizontally to increase dataset diversity.
 - Brightness, Contrast, and Color Perturbations: Altering these aspects to improve the model's robustness to lighting conditions.
 - Random Cropping: Extracting random patches from images to help the model focus on different parts of the object.
3. Test-Time Augmentation:
 - Generating multiple augmented versions of each test image and averaging the predictions to improve model performance.

Findings & Implications:

When comparing tissue samples with different stains, implementing these practices can be particularly beneficial:

- Centering and Normalization: Standardizing pixel values can help the model learn more effectively, especially when dealing with varying stain intensities.
- Data Augmentation: Applying techniques like random cropping and brightness adjustments can simulate variations in tissue samples, aiding the model in generalizing across different staining conditions.

Incorporating these strategies can enhance a model's ability to accurately assess similarities between tissue samples, leading to more reliable and robust performance.

Conclusions

1. Workarounds for Image Standardization

- **Padding Techniques:**

- Zero, Reflection, and Replication Padding: Common padding techniques, such as zero-padding, reflection, and replication, are used to standardize image sizes by adding pixels around the edges to achieve uniform dimensions. This method retains the central features and simplifies training by maintaining spatial consistency across samples.
- Trade-offs:
 - Positional Bias: As seen in [“Position, Padding, and Predictions,”](#) padding can introduce positional information, which might bias a model by making boundary areas distinctive in ways that could affect model generalization.
 - Boundary Artifacts: Padding, particularly zero-padding, may introduce boundary artifacts that affect convolution operations and may lead the model to focus on unrepresentative features.
- Effectiveness in Medical Imaging: Padding has been widely used in medical imaging, often as a straightforward solution to resize images. Reflection and replication padding, which smooth out edges without introducing zero-value pixels, may be more suitable in situations requiring spatial continuity.

- **Alternative Methods:**

- Hann Windows: The [PLOS article](#) on Hann windows for patch-based segmentation provides a unique approach that reduces edge effects by applying a window function. This technique, while primarily used in segmentation, may be adapted to reduce artifacts in medical image classification by smoothing pixel transitions, especially in large-scale images divided into patches.
- Learnable Resizing Modules: As demonstrated in the [Keras article](#), learnable resizing, where a model jointly learns resizing transformations with feature extraction, can effectively retain important details without padding, potentially reducing boundary artifacts. This method is useful in optimizing resizing specifically for the task and may yield higher classification accuracy for medical images, where each detail is crucial.
- Data Augmentation Techniques: The [Machine Learning Mastery article](#) underscores the effectiveness of techniques like cropping, rescaling, and random transformations. These can diversify the dataset and improve robustness to variations, helping the model generalize across samples with different orientations and sizes.

2. Trade-offs of Padding and Related Methods

- Information Loss vs. Consistency: Padding maintains image consistency across samples but can lead to boundary artifacts or information loss at the edges. Learnable resizing and data augmentation, though more computationally demanding, mitigate these issues by dynamically adapting the transformation for specific features.

- **Computational Efficiency:** Padding and simpler interpolation methods (like bilinear or bicubic resizing) are computationally inexpensive, making them practical for large datasets. However, the learnable resizing method, while computationally more intensive, can potentially enhance model performance by avoiding padding altogether.
- **Domain-Specific Application:** Techniques like Hann windows and reflection padding are especially relevant for medical imaging, as they retain spatial continuity and prevent boundary issues that could compromise the integrity of highly detailed images.

3. Use in Similar Problems and Effectiveness

- **Medical Imaging Studies:** Studies like the Hann window application and padding comparisons in CNNs show that careful selection of padding and resizing strategies can improve model performance in domains where details are crucial, such as medical image classification and segmentation.
- **Satellite and Other High-Detailed Image Domains:** In satellite imagery, the [Springer article on resizing techniques](#) discusses how these techniques directly influence classification confidence and accuracy. Bilinear interpolation, as shown in satellite image studies, balances accuracy and computational efficiency and could be equally beneficial for medical imaging where detail preservation is key.

Conclusion

For medical image classification, and our task at hand using eye tissue samples of different strains, padding offers a quick and consistent approach but may introduce unwanted artifacts and positional bias. Alternative techniques like Hann windows and learnable resizing modules appear to provide smoother boundary transitions and more optimized resizing, although perhaps at higher computational costs. By combining padding or windowing methods with data augmentation and task-specific resizing, you can mitigate the drawbacks of each method, improve image uniformity, and enhance classification accuracy, especially in detail-sensitive medical imaging.